

Spatio-Temporal Modeling and Simulation

Marcel Gietzmann-Sanders

2024

Dr. Andrew Seitz

Dr. Curry Cunningham

Michael Courtney, M.S.

College of Fisheries and Ocean Sciences
University of Alaska Fairbanks

Contents

| | | |
|----------|---|-----------|
| 1 | Log Odds Modeling | 3 |
| 1.1 | Fitting an Odds Model | 3 |
| 2 | Simulation | 5 |
| 2.1 | Dealing with Scale | 5 |
| 3 | Appendices | 10 |
| 3.1 | Convergence Issue with Non-Parametric Log Odds Models . . | 10 |
| 4 | Bibliography | 15 |

1 Log Odds Modeling

Objective 1. *To provide tooling that allows for using machine learning methods to fit models of the form*

$$\psi_k = G(\eta_k)$$

that maximize the following objective:

$$\mathcal{L} = \prod_i P'(v_i|\eta_i)$$

where:

$$P'(v_i|\eta_i) = \frac{e^{\psi_i}}{\sum_k e^{\psi_k}}$$

These models will be known as log odds models as they predict the "log odds for" each outcome v_k given the information contained in η_k .

1.1 Fitting an Odds Model

Our objective in using machine learning (ML) is to reduce the time spent looking for the particular form of $G(\eta_k)$ by instead allowing that form to be fit given the data at hand. Deep Learning (DL) models are especially well suited to this problem as they are both parametric (see Section 3.1 for why non-parametric models are not well suited to our problem) and have been shown to be able to represent just about any function given a large enough network [Benoit Lique and Nazarathy, 2024]. There is also a robust field around probabilistic deep learning that, as we shall see in a moment, we can take advantage of [Oliver Durr, 2020].

Taking a look at traditional probabilistic classifiers we can see that they have two things in common with our intended log odds modeling.

First, probabilistic classifiers use the categorical cross entropy cost function to optimize their weights. This cost function happens to be just another name for optimizing for the Log Likelihood of the data [Oliver Durr, 2020]. Therefore probabilistic DL shares the same objective function as we do.

Second, the final layer in a probabilistic DL network is the one that produces a probability for each class. This layer uses the softmax function as it's activation function which if $a_i = W_i x + b_i$ is given by $-\frac{e^{a_i}}{\sum_k e^{a_k}}$. This

means that if we can get our ψ_k to be the a_i that we'll be, in effect, training a log odds model.

We can do this by taking advantage of Keras' Functional API [Keras, 2023] which allows us to split our network into a branches and also share layers between branches. Specifically we can follow the following steps:

1. Define a series of layers that represent our underlying log odds model (these layers should end in a layer of output size 1 which will represent our ψ_k).
2. Split our training decisions into N choices.
3. Create a branch in our overall model for each of these N choices.
4. Converge the branches at a softmax layer where the weights are the identity matrix I and the biases all 0 (this ensures we directly pass through our ψ_k to the softmax function).
5. Use categorical cross entropy as our fitness function.

This architecture is illustrated by Fig. 1.

The main advantage of this architecture as opposed to a more classic approach to probabilistic DL is the fact that the weights are shared on each of the branches. This has the effect of drastically reducing the number of parameters in our network meaning that we can use far less data to train this model then we'd need to train a full bore, fully connected, probabilistic model. Each of our columns in Fig. 1 is in fact the same model and therefore updates coming from each of the columns goes to all of the columns.

This, then, solves one of our issues - reducing the variance in the model by switching to the log odds approach. But the other issue we're trying to solve is being able to predict on variable numbers of choices per decision. How does that work out here?

Well, in the case of training we know the maximum number of choices in any of our training decisions. Therefore we can set N to this maximum and for decisions where we have fewer than N choices we can simply provide some kind of default "do not choose" feature for the "missing" choices.

Later during inference on new data we can simply take our shared layers and use them to predict on each of the choices we're presented with. Just another demonstration that while we're using the architecture in Fig. 1 in order to *train* the log odds model, the actual log odds model is simply the shared layers in a single column behind the softmax layer.




Figure 1: Overall Architecture

2 Simulation

2.1 Dealing with Scale

2.1.1 Key Issues

There are a few key issues in trying to build spatio-temporal simulations using log odds models.

1. **Many World Problems.** *Issue: Exponentially growing state space. Solution: Grouping abundance by state keys.*

If we took the our log odds model outcomes at face value we'd have a simulation that quickly (and exponentially) careened out of control. Why is this? Well if we started with S distinct states and each of these generated C choices, in the next step we'd have SC states. Then in the next step SC^2 states. And so on until by step n we have SC^n states! Obviously this kind of exponential explosion of the state space is totally unmanageable. However we can get around this by recognizing that at the end of each step many of the states will overlap. For example if we are just predicting which new grid cells the fish are moving into, fish that start in nearby cells will move into the same grid cells as one another. Therefore we can group by the new resulting states, accumulate abundances from the various individual decisions, and keep our state space under control.

2. **Combinatorics of State.** *Issue: Problem scales by multiplying dimensions. Solution: Use cluster compute to take advantage of high parallelism in the problem.*

The size of each step in the simulation can get exceptionally large very quickly due to the number of dimensions that can be present in the state and the fact that these dimensions multiply on one another. So if I have 500 grid points, 50 distinct genetic cohorts, 10 different age groups, 3 size bins per age group, and 2 sexes that's a state space of $500 \bullet 50 \bullet 10 \bullet 3 \bullet 2 = 1,500,000$ distinct states to keep track of and model. And that would represent a relatively coarse description of the problem with no internal states present at all. Each additional dimension we add does not just add to this scale, it multiplies the scale. Very quickly then we will run into simulations that simply do not fit into a single machine's memory. However note that we can deal with each of these states independently which means that the parallelization opportunity is very high. Therefore if we take advantage of distributed compute paradigms like Spark we can parallelize the problems across as many machines as we like thereby preventing the bottleneck on single machine memory.

3. **Enormity of Results.** *Issue: Adding time and context increases the scale of the data even further. Solution: Using DBMS, optimized file storage systems, and a clear dimension vs fact delineation.*

In our example above we have 1.5 million entries per timestamp. If

we were interested in this over the course of a year and wanted to present results every 3 hours that would mean multiplying our 1.5 million states by approximately 3,000 different time points putting us at an easy 4.5 billion data points. Add into this the fact that useful exploration requires a lot of meta data (polygons, human readable timestamps, context like day/night, month of the year, country or state, etc.) and the sheer size of the data we would like to present and explore is quite daunting. We can do two things to make this data actually manageable. First is to take advantage of an actual database management system (DBMS). By writing there researchers can then interact with the data through queries, pull only what they need, and visualize just that. Then, in order to keep the data in simulation at a minimum we can separate our information into facts and dimensions [Ralph Kimball, 2013], facts being what is strictly needed for simulation and dimensions being repeated context that can be joined in later.

In summary then we want to ensure we obey the following design principles:

1. Collapse state at the end of each step.
2. Keep our simulation parallelizable by state and use cluster compute paradigms to allow for multi-machine parallelization.
3. Read from and write to databases or optimized file storage and retrieval systems.
4. Keep clear distinctions between facts and dimensions to reduce bloat in simulation while allowing for easy human exploration in database.

2.1.2 Kinds of Information

As we go through simulation there are a few key kinds of information each of which has specific implications for handling in simulation:

1. **Keys.** These are the information needed to identify a specific state. It includes spatial keys, temporal keys, demographic keys, and any internal state keys. We will parallelize along and group by these keys.
2. **Abundance.** This is a single metric that lets us know how much of an individual state there is.

3. **Joined Context.** This is any information that will be joined to the spatial or temporal keys and cannot be generated from the information already given. Examples would be things like bathymetry or temperatures. Because this information has to be joined on special consideration must be given for the sake of performance. If the data is small enough broadcast joins are best as they don't require reshuffling of the state information. Otherwise it is important that the states are partitioned by the join keys for this information so that as little shuffling as possible needs to happen during the joins.
4. **Built Context.** This is information that can be built from state keys or joined context. Examples would be sunset/sunrise times, season, etc. Whenever this kind of computation can be vectorized it should be for the sake of performance.
5. **Choice Independent Context.** This is information about the individual state itself that can be generated before choices are determined. This should obviously be computed before choices are generated.
6. **Choice Dependent Context.** This is any information specific to the choice in question and must be generated after choices have been built.
7. **Model Features.** This is information built from context that has to be normalized before being sent to the model.

From this we get a specific series of steps to follow:

1. Select or receive initial states (from DBMS).
2. Partition by state keys and join keys.
3. Join context.
4. Build choice independent context.
5. Build choices.
6. Build choice dependent context.

At this point it is important to know whether we are building data for training or running inference. In the former case (training):

1. Split on individuals input to create train and validation data.

2. Create and save normalization parameters from training data.
3. Build and normalize model features.
4. Capture choice selected.
5. Serialize and write to TFRecord (an optimized file storage for building deep learning models)

In the latter case (inference):

1. Build and normalize model features.
2. Apply the model to produce probabilities of selection.
3. Multiply probabilities by abundance.
4. Group by new states.
5. Write to DBMS.

As a final note it is important to recognize that the choices must be kept together until after model application in order to ensure the same decision stays on the same partition all the way until inference. This leads to very broad (but no less massive) datasets.

3 Appendices

3.1 Convergence Issue with Non-Parametric Log Odds Models

All standard Machine Learning (ML) pipelines assume that you have at least two things - targets and features. In our case we certainly have the latter but our target ψ_k is both unknown to us and also unmeasurable. So how are we to fit ML models if we have no target? In short, through iteration. Let's see how this can be done.

First some notation to help us. Our data is composed of a series of decisions $D_j = \{v_{jk}\}$ where j indicates each of the specific decisions and k the options within each decision. For each iteration we will build a model $\hat{G}_i(\eta_{jk})$ using the pairs $\psi_{jk(i-1)}, \eta_{jk}$. We will designate the outcome of that model ϕ_{jki} :

$$\phi_{jki} = \hat{G}_i(\eta_{jk})$$

Our probability is therefore:

$$P'_i(v_{jk}|\eta_{jk}) = \frac{e^{\phi_{jki}}}{\sum_p e^{\phi_{jpi}}}$$

Now if the k selected per decision D_j is given by s_j we want to maximize:

$$\mathcal{L} = \prod_j P'_i(v_{js_j}|\eta_{js_j}) \rightarrow \sum_j \ln [P'_i(v_{js_j}|\eta_{js_j})]$$

Using this information we will then propose an update u_{jki} s.t. $\psi_{jki} = \phi_{jki} + u_{jki}$ and repeat our iteration loop.

With that notation cleared up we can begin our iteration procedure.

First, let's assume we already have a guess for $\psi_{jk(i-1)}$. We can therefore train our model off of the ψ_{jk0}, η_{jk} pairs in standard ML fashion. Specifically we will fit a model that optimizes Mean Squared Error (the most common objective across ML software packages):

$$\min \left[\sum_{jk} \left(\hat{G}_i(\eta_{jk}) - \psi_{jk(i-1)} \right)^2 \right]$$

We now need to choose a set of updates u_{jki} . To get these we will turn to our overall objective function:

$$\sum_j \ln [P'_i(v_{js_j}|\eta_{js_j})] = \sum_j \ln \left[\frac{e^{\phi_{js_j i}}}{\sum_p e^{\phi_{jp i}}} \right]$$

Let's look at the gradient of this w.r.t the ϕ_{jki} . There are two cases:

$$\partial_{\phi_{js_j i}} \ln \mathcal{L} = \frac{1}{P'_i(v_{js_j}|\eta_{js_j})} \frac{\sum_p e^{\phi_{jp i}} - e^{\phi_{js_j i}}}{\left(\sum_p e^{\phi_{jp i}}\right)^2} e^{\phi_{js_j i}} = 1 - P'_i(v_{js_j}|\eta_{js_j})$$

$$\partial_{\phi_{j\neq j i}} \ln \mathcal{L} = \frac{1}{P'_i(v_{js_j}|\eta_{js_j})} \frac{-e^{\phi_{js_j i}}}{\left(\sum_p e^{\phi_{jp i}}\right)^2} e^{\phi_{j\neq j i}} = -P'_i(v_{j\neq j}|\eta_{j\neq j})$$

Next for point of illustration let's suppose that there are a set of ψ_{jki} which we'll designate as Z which share the same features η , i.e. our model has to give a single ϕ_{jki} for all such options. Our derivative then for that collection Z is given by:

$$\partial_Z \ln \mathcal{L} = \sum_{\phi_{js_j i} \in Z} (1 - P'_i(v_{js_j}|\eta_{js_j})) - \sum_{\phi_{j\neq j i} \in Z} P'_i(v_{j\neq j}|\eta_{j\neq j})$$

Given classic optimization tactics we know that our function will be maximized when these sums are 0 (and technically we'd also want to know that the second derivative was negative but we'll assume that's the case given how our iterations will work).

With this in mind let's now propose that our updates are given by:

$$u_{jki} = \alpha_i \partial_{\phi_{jki}} \ln \mathcal{L}$$

where α_i is a constant we'll call our "learning rate". Note that by using this update we will increase our ψ_{jki} guess where it corresponds to a taken option (s_j) and will decrease it where it corresponds to an option not taken ($\neq s_j$). This will therefore push us towards maximizing $\ln \mathcal{L}$ as opposed to minimizing it.

We are left with a final question - will our iteration sequence end when we've found the maximizing guesses of ψ_{jki} ? To answer this we turn back to the term we are maximizing when fitting the \hat{G}_i :

$$\min \left[\sum_{jk} \left(\hat{G}_i(\eta_{jk}) - \psi_{jk(i-1)} \right)^2 \right]$$

Our new fit will look like:

$$\sum_{jk} \left(\hat{G}_{i+1}(\eta_{jk}) - (\hat{G}_i(\eta_{jk}) + u_{jki}) \right)^2 = \sum_{jk} \left(\delta \hat{G}_{jk} - u_{jki} \right)^2$$

Given our Z once again we can take the derivative w.r.t $\delta \hat{G}_{jk}$ where $\phi_{jki} \in Z$.

$$\partial_Z \left[\sum_{jk} \left(\delta \hat{G} - u_{jki} \right)^2 \right] = \sum_{\phi_{jki} \in Z} 2 \left(\delta \hat{G} - u_{jki} \right) = 2 \sum_{\phi_{jki} \in Z} \delta \hat{G} - 2 \sum_{\phi_{jki} \in Z} u_{jki}$$

But now remember that if we've maximized w.r.t Z that:

$$\sum_{\phi_{jki} \in Z} u_{jki} = 0$$

which means that in order for our partial derivative above to be zero (and therefore our error term be at a minimum) that $\delta \hat{G} = 0$. And this means our iteration will have stopped!

Procedure 1. *Fitting a Log Odds Model*

1. Collect decisions D_j and corresponding features η_{jk} , options v_{jk} , and selection s_j .
2. Make an initial guess $\psi_{jk0} = 0$.
3. Fit \hat{G}_i on the pairs of $\psi_{jk(i-1)}, \eta_{jk}$ using MSE to produce the ϕ_{jki} .
4. Generate the $u_{jki} = \alpha_i \partial_{\phi_{jki}} \ln \mathcal{L}$ and produce a new set of ψ_{jki} .
5. Repeat 3 and 4 until, varying α_i until convergence.

$$\partial_{\phi_{js_j i}} \ln \mathcal{L} = 1 - P'_i(v_{js_j} | \eta_{js_j})$$

$$\partial_{\phi_{j\not{s}_j i}} \ln \mathcal{L} = -P'_i(v_{j\not{s}_j} | \eta_{j\not{s}_j})$$

There is however an issue with this approach. Recall that our gradient is:

$$\partial_Z \ln \mathcal{L} = \sum_{\phi_{js_j i} \in Z} (1 - P'_i(v_{js_j} | \eta_{js_j})) - \sum_{\phi_{j\not{s}_j i} \in Z} P'_i(v_{j\not{s}_j} | \eta_{j\not{s}_j})$$

Let's however look at the second derivative of $\ln \mathcal{L}$:

$$\partial_{\phi_{jk}} P'_i(v_{jk} | \eta_{jk}) = \partial_{\phi_{jk}} \frac{e^{\phi_{jk}}}{\sum_p e^{\phi_{jp}}} = (1 - P'_i(v_{jk} | \eta_{jk})) P'_i(v_{jk} | \eta_{jk})$$

$$\partial_Z^2 \ln \mathcal{L} = - \sum_{\phi_{js_j i} \in Z} (1 - P'_i(v_{js_j} | \eta_{js_j})) P'_i(v_{js_j} | \eta_{js_j}) - \sum_{\phi_{j\not{s}_j i} \in Z} (1 - P'_i(v_{j\not{s}_j} | \eta_{j\not{s}_j})) P'_i(v_{j\not{s}_j} | \eta_{j\not{s}_j})$$

What's important to note is that this function will be near 0 at $P \approx 1$ or $P \approx 0$ and will have its largest magnitude near $P \approx 0.5$. Why is this an issue? Well we know that as we get close to our maximum value for $\ln \mathcal{L}$ that our steps towards that maximum will get smaller and smaller. This is just because our steps are based on a derivative and we are seeking where the derivative is zero. Put another way getting to our max gets harder the closer we get to that maximum.

Now normally one deals with this by using the second derivative (the curvature) as a kind of correction. As you get closer to your maximum, you

can use your curvature to guide how quickly you can move. If the curvature is very small in magnitude you can move more quickly because the odds of you overshooting your maximum are smaller. If your curvature is very large you'll slow things down because you know only small steps are required to make big changes to the derivative (the thing we are ultimately trying to set to a specific value here). I.e. you can modulate your step size by the inverse of the curvature.

However, in our case we don't actually know what the curvature is because we don't know what Z is! Therefore we have to plan for the most volatile case - the case where $P \approx 0.5$. And that means that for any of our Z 's where P is approaching 1 or 0 we'll be moving at a snail's pace (convergence will take forever).

Now in case you're wondering whether we could solve the $P \approx 0.5$ case first and then move onto the others remember that anytime we change any of the ϕ all the other probabilities change. So we have to solve this problem all at once. Because of that and because we cannot depend on knowing the Z (different Z can give the same value ϕ) we're stuffed with having to take extraordinarily long convergence times. And this more or less means we've got no shot of using this in practice.

Log odds modeling doesn't work in practice for non-parametric models.

4 Bibliography

References

- [Benoit Lique and Nazarathy, 2024] Benoit Lique, S. M. and Nazarathy, Y. (2024). *Mathematical Engineering of Deep Learning*. CRC Press.
- [Keras, 2023] Keras (2023). Keras functional api.
- [Oliver Durr, 2020] Oliver Durr, Beate Sick, E. M. (2020). *Probabilistic Deep Learning*. Manning Publications.
- [Ralph Kimball, 2013] Ralph Kimball, M. R. (2013). The data warehouse toolkit: The definitive guide to dimensional modeling.